

Prompting lab

Five paired prompts for the OpenAI Playground. Try the **weak** version first, then the **strong** one. Same model, same facts. Only the brief changes. Tap Copy to paste into the prompt box.

How to use this: Open platform.openai.com/chat, copy a prompt below, and paste it in. Each weak/strong pair shares the same source material so you can compare fairly. For pair 2, set Text format to json_schema. For pair 4, try changing Reasoning effort.

Suggested order:

1 Specificity 3 Few-shot 2 Structured output 4 Chain of thought 5 Constraints

PAIR 1

Specificity and role framing

What it is: Tell the model who it is, who it writes for, how long to write, what tone to use, and which facts must appear. Vague asks for five guesses at once.

WEAK

Copy weak

Write about our new feature.

Product notes:

- Product: B2B analytics tool
- Feature: anomaly detection
- Flags outliers within 60 seconds
- Requires no configuration
- Audience we care about: data engineers evaluating tools

STRONG

Copy strong

You are a product marketing writer for a B2B analytics tool. Write a 120-word announcement for our new anomaly detection feature. Audience: data engineers evaluating tools. Tone: direct, no hype. Mention that it flags outliers within 60 seconds and requires no configuration.

Why it works: both versions have the same product facts. The strong version adds role, length, audience, and tone. The weak version still forces the model to guess how to use those facts.

PAIR 2

Structured output

What it is: Ask for a fixed schema (JSON fields and types) instead of free prose. In Playground, set Text format to json_schema so the API enforces the shape.

WEAK

[Copy weak](#)

Pull out the important stuff from this support ticket.

Ticket: "Hi, I've been on the Pro plan since March and the export button has done nothing for three days. I've got a client report due Friday. Please help."

STRONG

[Copy strong](#)

Extract the following from the support ticket below and return valid JSON only, no prose or markdown fences.

Schema:

```
{"customer_tier": "free|pro|enterprise", "issue_category":  
"billing|bug|feature_request|other", "severity": 1-5, "summary": "string, max 20  
words", "action_required": true|false}
```

If a field cannot be determined, use null.

Ticket: "Hi, I've been on the Pro plan since March and the export button has done nothing for three days. I've got a client report due Friday. Please help."

Why it works: same ticket in both. Set Text format to json_schema in the right panel. The strong

brief names the fields; the weak one leaves "important" undefined.

PAIR 3

Few-shot prompting

What it is: Show a few complete input → output examples, then leave the next input unfinished. The examples teach labels, format, and how detailed the reason should be.

WEAK

[Copy weak](#)

Classify the sentiment of these product reviews.

Review: "Setup took ten minutes and it just worked."

Review: "Does what it says, though the mobile app lags."

Review: "Third outage this month. We're migrating."

Review: "Support answered in two minutes but couldn't fix it."

STRONG

[Copy strong](#)

Classify each review's sentiment. Follow the format of the examples exactly.

Review: "Setup took ten minutes and it just worked."

Sentiment: positive

Reason: effortless onboarding

Review: "Does what it says, though the mobile app lags."

Sentiment: mixed

Reason: core function fine, performance issue

Review: "Third outage this month. We're migrating."

Sentiment: negative

Reason: reliability failure, churn signal

Review: "Support answered in two minutes but couldn't fix it."

Sentiment:

Why it works: both versions get the same four reviews. The strong version adds labeled examples so labels, reason depth, and output shape are locked in. The weak version leaves all of that to

guesswork.

PAIR 4

Chain of thought

What it is: Ask the model to work step by step and put the final number on a fixed line. You can audit the path, not only the answer.

WEAK

Copy weak

A subscription costs £40/month per seat, with a 15% discount if the year is paid upfront. A team of 12 pays annually upfront in month 1. In month 7 they add 5 seats, billed monthly at full price for the remaining months of the year. What is the total first-year cost? Just give me the number.

STRONG

Copy strong

Work through this step by step. Show your reasoning, then give the final answer on a line starting with "ANSWER:".

A subscription costs £40/month per seat, with a 15% discount if the year is paid upfront. A team of 12 pays annually upfront in month 1. In month 7 they add 5 seats, billed monthly at full price for the remaining months of the year. What is the total first-year cost?

Why it works: same numbers in both. Try Reasoning effort medium then high on the weak version. More effort can help, but asking for steps plus a fixed ANSWER: line still makes the result easy to check.

PAIR 5

Constraints and negative instructions

What it is: Replace soft rules ("not too long", "no jargon") with checkable limits: counts, word caps, and what to do when information is missing.

WEAK

Copy weak

Summarise this research paper but don't make it too long and don't use jargon.

Paper: "We surveyed 412 software teams adopting coding agents over six months. Teams that wrote explicit done-conditions before generation shipped 28% fewer reverted PRs than teams that prompted ad hoc. Limitation: all teams were already using GitHub Actions, so results may not generalise to teams without CI."

STRONG

Copy strong

Summarise the paper below for a non-technical executive.

Rules:

- Exactly three bullet points
- Each bullet under 25 words
- No acronyms unless expanded on first use
- State one limitation of the study
- If the paper does not report a limitation, write "No limitation stated"

Paper: "We surveyed 412 software teams adopting coding agents over six months. Teams that wrote explicit done-conditions before generation shipped 28% fewer reverted PRs than teams that prompted ad hoc. Limitation: all teams were already using GitHub Actions, so results may not generalise to teams without CI."

Why it works: same paper in both. Soft rules like "not too long" are vague. "Exactly three bullets under 25 words" is something you can check at a glance.

Weak first. Strong second. Same model, different brief.

That gap is prompt engineering — and the schemas you practise here map cleanly into typed API responses later in the bootcamp.